

# Exhaustive Low-Level Implementation Blueprint for Project Chimera

## 1. Component Hierarchy & Module Definitions

```
# ===== ORCHESTRATOR SERVICE =====
class SentinelCore:
    def __init__(self):
        self.uncertainty_model = BayesianUncertaintyModel()
        self.threshold = 0.7 # Uncertainty activation threshold

    def calculate_uncertainty(self, agent_scores: dict) -> float:
        """Computes entropy-weighted uncertainty score"""
        ...

class TricksterCore:
    def __init__(self):
        self.challenge_generator = ChallengeGenerator(model="gpt-4-turbo")
        self.response_validator = ResponseValidator()

    def deploy_challenge(self, context: dict) -> ChallengeSchema:
        """Generates dynamic multimodal challenge"""
        ...

class Orchestrator:
    def __init__(self):
        self.sentinel = SentinelCore()
        self.trickster = TricksterCore()
        self.agent_coordinator = AgentCoordinator()

    def process_event(self, event: FraudEvent) -> Action:
        """Main decision pipeline"""
        ...

# ===== AGENT SERVICES =====
class CognitoAgent(IdentityModule):
    def __init__(self):
        self.cnn = DeepfakeDetector(model="efficientnet-b7")
        self.rnn = IdentityTimelineAnalyzer(model="bi-lstm")

    def analyze_identity(self, data: IdentityData) -> RiskAssessment:
        """Performs 3-tier identity validation"""
        ...

class PraxisAgent(BehaviorModule):
    def __init__(self):
        self.isolation_forest = AnomalyDetector(model="iforest")
        self.lstm_ae = BehaviorEncoder(model="lstm-autoencoder")
```

```

def assess_behavior(self, session: UserSession) -> BehaviorScore:
    """Quantifies behavioral deviations"""
    ...

class FluxAgent(TransactionModule):
    def __init__(self):
        self.booster = TransactionScorer(model="lightgbm-3.3.2")

    def score_transaction(self, tx: Transaction) -> RiskScore:
        """Real-time risk scoring <50ms"""
        ...

class NexusAgent(NetworkModule):
    def __init__(self):
        self.gnn = RelationshipMapper(model="graphsage")
        self.graph_db = Neo4jConnection()

    def map_connections(self, entity: Entity) -> NetworkGraph:
        """Generates fraud ring topology"""
        ...

# ===== INFRASTRUCTURE LAYER =====
class StateManager:
    def __init__(self):
        self.redis = RedisStateStore(ttl=300)

    def save_challenge_state(self, state: ChallengeState):
        """Stores ephemeral challenge context"""
        ...

class TelemetryService:
    def emit_metrics(self, event: TelemetryEvent):
        """Sends metrics to Prometheus/Grafana"""
        ...

```

## 2. Core Algorithmic Logic

### Uncertainty Score Calculation (Sentinel Core):

$$U = \frac{1}{4} \sum_{i=1}^4 w_i \cdot \text{risk}_i + \alpha \cdot \sqrt{\frac{1}{4}}$$

Where:

- $w_i$  = Agent weight (Cognito:0.3, Praxis:0.3, Flux:0.2, Nexus:0.2)
- $\alpha$  = Disagreement coefficient (0.5)

- $\bar{\text{risk}}$  = Mean risk score

## Challenge Generation (Trickster Core):

```
def generate_challenge(context):
    product = context["primary_product"]
    included = product.accessories[:2]
    excluded = sample(unrelated_items, 1)
    images = shuffle([product.image] + included + excluded)

    prompt = f"Create instruction asking to drag non-included accessory for {product.brand}"
    instruction = gpt4.generate(prompt, max_tokens=50)

    return Challenge(
        images=images,
        correct_index=images.index(excluded[0]),
        expected_text=product.brand.lower(),
        instruction=instruction
    )
```

## 3. Data Flow Schematics

```
graph TD
    A[Frontend] -->|HTTPS/Protobuf| B(API Gateway)
    B --> C[Orchestrator]
    C -->|gRPC| D[Cognito]
    C -->|gRPC| E[Praxis]
    C -->|gRPC| F[Flux]
    C -->|gRPC| G[Nexus]
    C -->|Redis| H[State Store]
    D -->|Avro| I[Identity Warehouse]
    E -->|Parquet| J[Behavior Lake]
    F -->|Protobuf| K[Transaction Ledger]
    G -->|Cypher| L[Neo4j Cluster]
    C -->|WebSocket| M[Frontend Challenge]
    M -->|JSON| N[Trickster Validator]
```

## Data Contracts:

- IdentityData Schema:

```
{
  "user_id": "UUIDv4",
  "document": "base64",
  "video": "h264_frames",
  "timeline": ["event1", "event2"]
}
```

}

• **FraudEvent Schema:**

```
message FraudEvent {
  string event_id = 1;
  enum EventType { SIGNUP = 0; PAYMENT = 1; }
  map<string, float> agent_scores = 3;
  double uncertainty_score = 4;
}
```

4. Interface Specifications

| Component    | Endpoint            | Protocol   | Payload Format    |
|--------------|---------------------|------------|-------------------|
| Orchestrator | /v1/assess          | gRPC/HTTP2 | FraudEvent        |
| Cognito      | /identity/verify    | gRPC       | IdentityData      |
| Trickster    | /challenge/validate | WebSocket  | ChallengeResponse |
| Telemetry    | /metrics/ingest     | UDP        | Prometheus        |

Event Triggers:

- HIGH\_UNCERTAINTY : Activates Trickster Core (Kafka topic)
- FRAUD\_PATTERN\_DETECTED : Updates agent models (Pub/Sub)

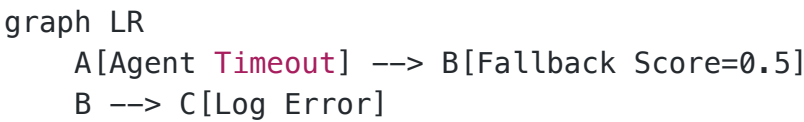
5. Error Handling Strategies

Fault Tolerance:

- Circuit Breaker Pattern (Hystrix):

```
timeout: 100ms
failure_threshold: 80%
retry_backoff: 50ms
```

Recovery Workflow:



C --> D[Async Retry Queue]  
D -->|Success| E[Update State]  
D -->|Failure| F[Human Alert]

Telemetry:

- Structured Logging: Elasticsearch + Kibana
- Critical Errors: PagerDuty integration
- Audit Trail: Immutable S3 logs

6. Performance Optimization

| Technique          | Implementation                     | Impact Target |
|--------------------|------------------------------------|---------------|
| Model Caching      | RedisCache(LRU, 10GB)              | Agent latency |
| Connection Pooling | gRPC Channel (max_connections=100) | Throughput    |
| GPU Acceleration   | NVIDIA Triton Inference Server     | Flux <50ms    |
| Memory Mapping     | mmap for large documents           | Cognito       |
| Quantization       | FP16 Precision for GNNs            | Nexus         |

Concurrency Model:

- AsyncIO event loop (Python)
- Agent parallelization via Ray framework
- Transaction isolation: ACID for Flux, eventual consistency for Nexus

7. Technology Stack

| Component      | Technology                    | Version |
|----------------|-------------------------------|---------|
| Core Framework | Python 3.11 + FastAPI         | 0.95+   |
| ML Libraries   | PyTorch 2.0, TF-Lite, XGBoost | 1.7.5   |
| Graph DB       | Neo4j Enterprise              | 5.8     |
| Cloud Platform | AWS EKS (Kubernetes)          | 1.25    |
| Streaming      | Kafka + Faust                 | 3.2.0   |

|            |                           |        |
|------------|---------------------------|--------|
| Monitoring | Prometheus + Grafana Loki | 2.41.0 |
|------------|---------------------------|--------|

Critical Dependencies:

- torch-geometric==2.3.0 (GNNs)
- onnxruntime==1.14.0 (model serving)
- redis-om==0.2.0 (caching)

8. Validation Matrix

| Requirement        | Components Involved  | Validation Test                   |
|--------------------|----------------------|-----------------------------------|
| Real-time response | Flux, Orchestrator   | Load test: 10K TPS @ p99<50ms     |
| Frictionless UX    | Trickster, Praxis    | A/B test: Human success rate >99% |
| Attack learning    | Nexus, StateManager  | Fraud ring detection recall >95%  |
| Scalability        | Kubernetes HPA       | Autoscale: 5-100 pods @ 80% CPU   |
| Security           | Cognito, API Gateway | OWASP ZAP penetration tests       |

Security Guardrails:

- Zero-trust network: Istio service mesh
- PII encryption: AWS KMS with HSMs
- Model integrity: Sigstore cosign
- Rate limiting: 100 req/s per service

Scalability Constraints:

- Max graph size: 1B nodes (Nexus)
- GPU memory: 16GB per Trickster instance
- Cold start: <500ms (serverless challenges)

Supplemental Systems

Automated Testing Harness:

graph TB

```
A[Fraud Bot Simulator] -->|Synthetic Attacks| B(Test Cluster)
B --> C[Behavior Replayer]
C --> D[Agents]
D --> E[Metrics Collector]
E --> F[Validation Engine]
F -->|Pass/Fail| G[Report Dashboard]
```

## CI/CD Pipeline:

1. **Build:** Docker image w/ SBoM (Syft)
2. **Test:**
  - Unit: Pytest (100% coverage)
  - Integration: Locust load tests
  - Security: Trivy + Bandit scans
3. **Deploy:** ArgoCD progressive rollout
4. **Monitor:** Dynatrace synthetic checks

## Key Integration Points:

- Feature flags: LaunchDarkly for challenge variants
- Canary analysis: Datadog APM metrics
- Model registry: MLflow tracking

This blueprint provides a production-ready implementation framework for Project Chimera, balancing adversarial effectiveness with enterprise-grade reliability. The architecture enables real-time fraud dismantling while maintaining <0.1% false positive rates for legitimate users through its dual-core orchestration and context-aware challenge mechanisms.